
Quiz: Retrieval-Augmented Generation (RAG)

Name: _____ Himani Singh _____

Date: _____ 16 March 2026 _____

Score: _____ / 20

Section A: Multiple Choice Questions (1 mark each)

1. What is the main purpose of Retrieval-Augmented Generation (RAG)?

- A. To compress large datasets
 - B. To combine document retrieval with language generation
 - C. To train a neural network from scratch
 - D. To replace vector databases
-

2. In a RAG system, what is the role of a vector database?

- A. Store trained models
 - B. Store embeddings for similarity search
 - C. Generate answers
 - D. Train transformers
-

3. What does “chunking” refer to in RAG pipelines?

- A. Dividing a document into smaller sections
 - B. Encrypting documents
 - C. Compressing text data
 - D. Converting text into images
-

4. Which method is commonly used to convert text chunks into vectors?

- A. Tokenization

B. Embeddings

C. Sorting

D. Indexing

5. Which component generates the final response in a RAG system?

A. Vector database

B. Embedding model

C. Large Language Model

D. Data loader

Section B: True or False (1 mark each)

6. RAG helps LLMs use external knowledge sources.

True / False

7. Chunking large documents improves retrieval accuracy.

True / False

8. Context injection means adding retrieved information into the prompt.

True / False

9. In RAG, the LLM directly searches PDFs without embeddings.

True / False

10. Evaluating a RAG system only requires checking model accuracy.

True / False

Section C: Short Answer (2 marks each)

11. Write the basic workflow of a RAG system.

ANSWER: Document ingestion → Chunking → Create embeddings → Store in vector database → User query → Retrieve relevant chunks → Add context to prompt → LLM generates answer.

12. Why is chunk size important in RAG pipelines?

ANSWER: Chunk size affects retrieval accuracy. If chunks are too large, irrelevant information may be retrieved; if too small, important context may be lost.

13. What is context injection and why is it useful?

ANSWER: Context injection is the process of adding retrieved document information into the prompt given to the LLM so it can generate answers based on relevant external knowledge.

14. Name two evaluation metrics or considerations for RAG systems.

ANSWER: Retrieval accuracy / relevance of retrieved documents

Answer correctness / response quality

15. Give one real-world application of RAG.

ANSWER: Document question answering systems, such as legal research assistants that retrieve information from regulations and generate explanations.

✓ Total Marks: 20

-